

GOVERNED MULTI-AGENT WORKFLOW

Product safety evidence checking, governed from prompt to identity.

RecallGuard AI checks vendor product submissions against product safety policy, recall evidence, and KC certification records before marketplace listing.

KNOWLEDGE AGENT

TASK AGENT

TRACES + ENTRA ID

2

AGENTS

4

TEST CASES

1

WORKFLOW

WHY GOVERNANCE MATTERS

A vendor CSV can carry more risk than it looks.

The workflow has to catch product risk, missing evidence, and adversarial content without inventing a compliance answer.

01

Recalled product slips through

A recalled model or importer match must create HOLD and trigger human review.

02

Certification evidence is incomplete

Missing model, manufacturer, or KC number should become REVIEW, not APPROVE.

03

Vendor note attempts instruction override

Uploaded files are untrusted data; notes must not change agent policy.

AUTOGRADER-CLEAR COVERAGE

Every final-activity requirement maps to a concrete RecallGuard artifact.

REQUIREMENT	BUILT ARTIFACT	FOUNDRY EVIDENCE	STATUS
Knowledge Agent	recallguard-knowledge-agent	File Search + vector store	COMPLETE
Task Agent	recallguard-task-agent-v7-public-data	Code Interpreter + checker	COMPLETE
Sequential workflow	recallguard-governed-workflow-v5-public-data	Workflow agent + portal runbook	COMPLETE
Traces + Preview	Success, missing, edge runs	Live API artifacts + portal visuals	COMPLETE
Guardrails + identity	Prompt-injection controls + Entra IDs	Agent principal IDs documented	COMPLETE

ORCHESTRATION

Ground first. Act second. Escalate when risk is real.

The workflow intentionally separates policy grounding from file execution, then preserves a human gate for HOLD decisions.



OBSERVABILITY RAIL

Preview + Traces + Guardrails + Entra Agent ID governance sit across the full workflow.

GROUNDING Q&A

The policy answer is allowed to say “I do not know.”

The Knowledge Agent uses indexed policy files and must cite evidence or surface missing information instead of guessing.

INSTRUCTION SNIPPET


Answer only using connected knowledge sources. Do not guess. List evidence, missing information, and next compliance action.

KNOWLEDGE SOURCE	ROLE IN ANSWERING
product_safety_review_sop.md	Decision labels and conservative approval rule
kc_certification_review_checklist.md	Required vendor fields and evidence strength
recall_response_policy.md	Recall match handling and human review rule
vendor_submission_requirements.md	Untrusted content and required submission fields

TOOL EXECUTION

A false HOLD became the reason to stop trusting free-form classification.

The final Task Agent uses Code Interpreter plus recallguard_checker.py so row-level decisions are deterministic and evidence_type-aware.



EVIDENCE MATRIX

The demo covers the happy path, missing evidence, recall risk, and prompt injection.

1

HOLD

2

REVIEW

These outcomes are intentionally submission-friendly: every result maps to a Foundry final-activity requirement.

TEST	INPUT	OBSERVED OUTCOME	WHY IT MATTERS
Knowledge Q&A	Policy question	Grounded answer	Covers Knowledge Agent requirement
Recall match	vendor_products_recall_match.csv	1 HOLD / 1 APPROVE	Proves risk triage + HITL trigger
Missing fields	vendor_products_missing_fields.csv	2 REVIEW	Prevents unsafe approval
Prompt injection edge	vendor_products_prompt_injection.csv	1 HOLD / 1 APPROVE	Treats notes as untrusted data

GUARDRAILS + IDENTITY

The workflow is governed at the prompt, tool, trace, and identity layers.


Foundry guardrails cover prompt injection and safety. Agent instructions and deterministic code add application-level controls.

CONTROL	IMPLEMENTATION	EVIDENCE	KEY IDS
Prompt injection	Vendor notes treated only as data	Edge case output saved	Knowledge 934a3b63...
Ungrounded answers	Knowledge Agent cites connected sources	File Search/vector store	Task v7 0173cc3c...
Unsafe approval	Missing identifiers ⇒ REVIEW	Checker policy	Workflow 87a0ae4f...
Identity governance	Agent principal IDs documented	Entra Agent ID notes	

3-MINUTE DEMO SCRIPT

The recording should prove Foundry usage before it explains the idea.


Use this slide as the video run-of-show. Replace the placeholder segments with actual portal Preview, Trace, Guardrails, and Agent ID screens.



READY-TO-SUBMIT FILES

Submit the video and report; keep the PPT as the editable source.

The report is ATS-friendly, the deck is recording-ready, and the MP4 can be regenerated from slide previews and narration.



FILE	PURPOSE	STATUS
RecallGuard_AI_Demo_Somi.mp4	Required video deliverable	READY
RecallGuard_AI_Build_Report_Somi.pdf	Required document deliverable	READY
RecallGuard_AI_Foundry_Deck_Somi.pptx	Editable video/source deck	READY
RecallGuard_AI_PPT_Video_Script_Somi.md	Narration and capture guide	READY